

RWorksheet_Moquete#4c.Rmd

Renz Moquete

2025-12-12

1. Use the dataset mpg

a. Show your solutions on how to import a csv file into the environment.

Since mpg is already available in ggplot2, we'll load it directly

```
data(mpg)

# To demonstrate importing a CSV file (if we had one):
# mpg_csv <- read_csv("mpg.csv") # For CSV file
# mpg_excel <- read_excel("mpg.xlsx") # For Excel file

cat("First 6 rows of mpg dataset:\n")

## First 6 rows of mpg dataset:
print(head(mpg))

## # A tibble: 6 x 11
##   manufacturer model displ year cyl trans     drv   cty   hwy fl class
##   <chr>        <chr> <dbl> <int> <int> <chr>   <chr> <int> <int> <chr> <chr>
## 1 audi         a4      1.8  1999    4 auto(15) f       18    29 p     compa-
## 2 audi         a4      1.8  1999    4 manual(m5) f      21    29 p     compa-
## 3 audi         a4      2.0  2008    4 manual(m6) f      20    31 p     compa-
## 4 audi         a4      2.0  2008    4 auto(av)   f      21    30 p     compa-
## 5 audi         a4      2.8  1999    6 auto(15)  f      16    26 p     compa-
## 6 audi         a4      2.8  1999    6 manual(m5) f      18    26 p     compa-
```

b. Which variables from mpg dataset are categorical?

```
cat("\nCategorical variables in mpg dataset:\n")

##
## Categorical variables in mpg dataset:
categorical_vars <- names(mpg)[sapply(mpg, is.character)]
print(categorical_vars)

## [1] "manufacturer" "model"          "trans"           "drv"            "fl"
## [6] "class"

# Also check factor variables
cat("\nFactor variables in mpg dataset:\n")

##
## Factor variables in mpg dataset:
```

```

factor_vars <- names(mpg)[sapply(mpg, is.factor)]
print(factor_vars)

## character(0)

```

c. Which are continuous variables?

```

cat("\nContinuous/numeric variables in mpg dataset:\n")

##
## Continuous/numeric variables in mpg dataset:
numeric_vars <- names(mpg)[sapply(mpg, is.numeric)]
print(numeric_vars)

## [1] "displ" "year"  "cyl"   "cty"   "hwy"

```

2. Which manufacturer has the most models in this data set? Which model has the most variations?

- a. Group the manufacturers and find the unique models

```

manufacturer_summary <- mpg %>%
  group_by(manufacturer) %>%
  summarise(
    num_models = n_distinct(model),
    total_cars = n()
  ) %>%
  arrange(desc(num_models))

cat("Manufacturer with most models:\n")

## Manufacturer with most models:
print(manufacturer_summary[1, ])

## # A tibble: 1 x 3
##   manufacturer num_models total_cars
##   <chr>           <int>      <int>
## 1 toyota            6          34

# Model with most variations
model_summary <- mpg %>%
  group_by(model) %>%
  summarise(
    num_variations = n(),
    manufacturers = paste(unique(manufacturer), collapse = ", ")
  ) %>%
  arrange(desc(num_variations))

cat("\nModel with most variations:\n")

##
## Model with most variations:
print(model_summary[1, ])

## # A tibble: 1 x 3

```

```

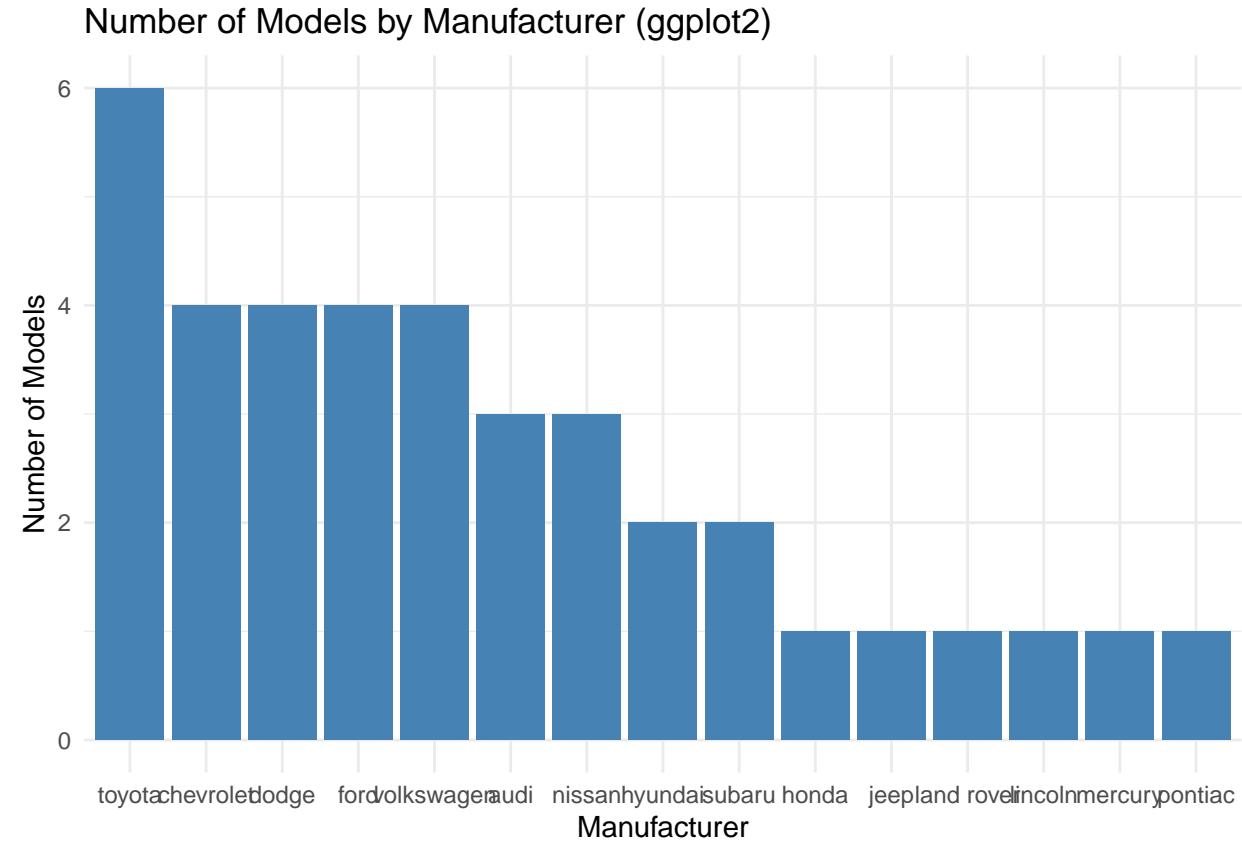
##   model      num_variations manufacturers
##   <chr>          <int> <chr>
## 1 caravan 2wd           11 dodge

b. Graph the result by using plot() and ggplot()

# Base R plot
par(mfrow = c(1, 2))
barplot(manufacturer_summary$num_models,
        names.arg = manufacturer_summary$manufacturer,
        main = "Number of Models by Manufacturer (Base R)",
        xlab = "Manufacturer",
        ylab = "Number of Models",
        col = "lightblue",
        las = 2,
        cex.names = 0.7)

# ggplot2
ggplot(manufacturer_summary, aes(x = reorder(manufacturer, -num_models), y = num_models)) +
  geom_bar(stat = "identity", fill = "steelblue") +
  labs(title = "Number of Models by Manufacturer (ggplot2)",
       x = "Manufacturer",
       y = "Number of Models") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1)) +
  theme_minimal()
par(mfrow = c(1, 1))

```



c. Same dataset will be used. Show relationship of model and manufacturer

```

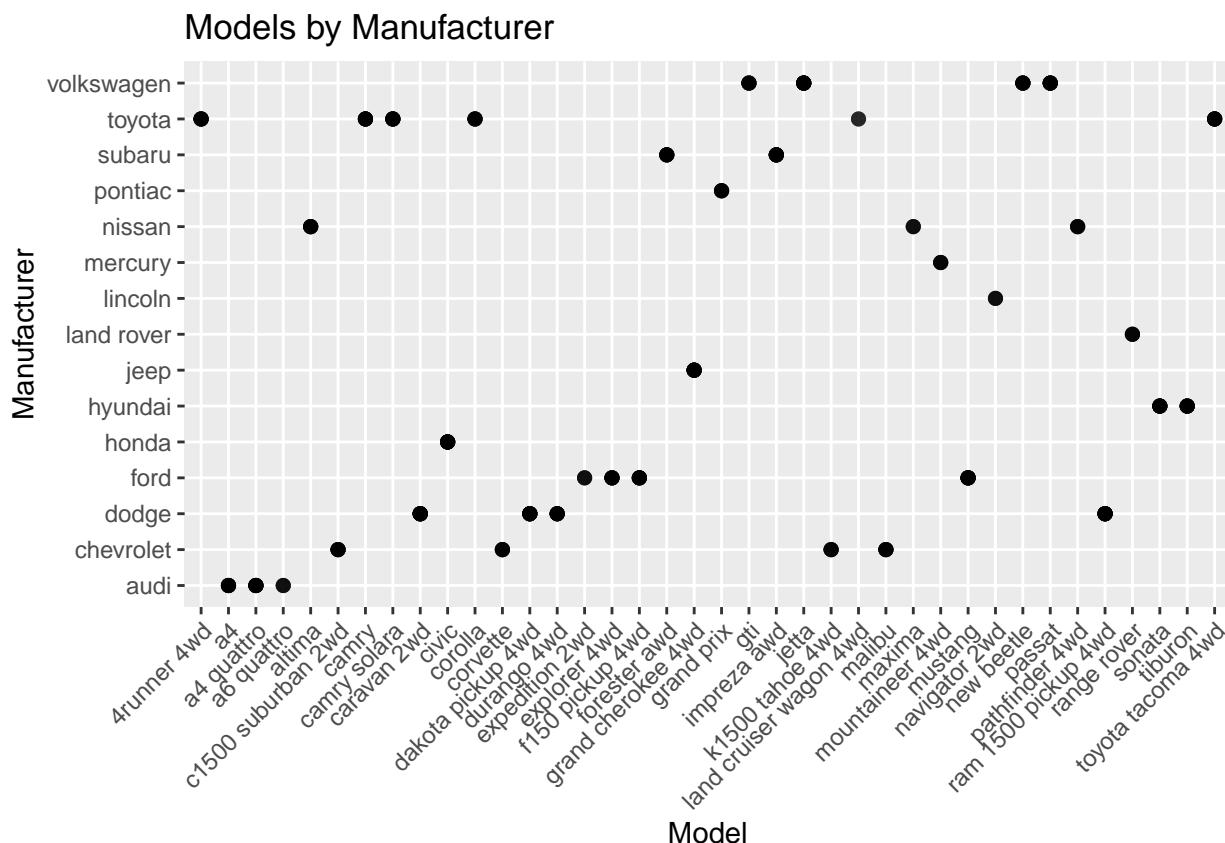
# a. What does ggplot(mpg, aes(model, manufacturer)) + geom_point() show?
cat("This plot shows each car model as a point, positioned by model (x-axis) and manufacturer (y-axis).")

## This plot shows each car model as a point, positioned by model (x-axis) and manufacturer (y-axis).
cat("It visualizes which manufacturers produce which models.\n\n")

## It visualizes which manufacturers produce which models.

# Create the plot
p1 <- ggplot(mpg, aes(x = model, y = manufacturer)) +
  geom_point(alpha = 0.6, size = 2) +
  labs(title = "Models by Manufacturer",
       x = "Model",
       y = "Manufacturer") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
print(p1)

```



```

# b. For you, is it useful? If not, how could you modify the data to make it more informative?
cat("\nb. This plot could be improved by:\n")

```

```

##
## b. This plot could be improved by:
cat("1. Adding color to show year or other variables\n")

## 1. Adding color to show year or other variables
cat("2. Adding jitter to reduce overplotting\n")

```

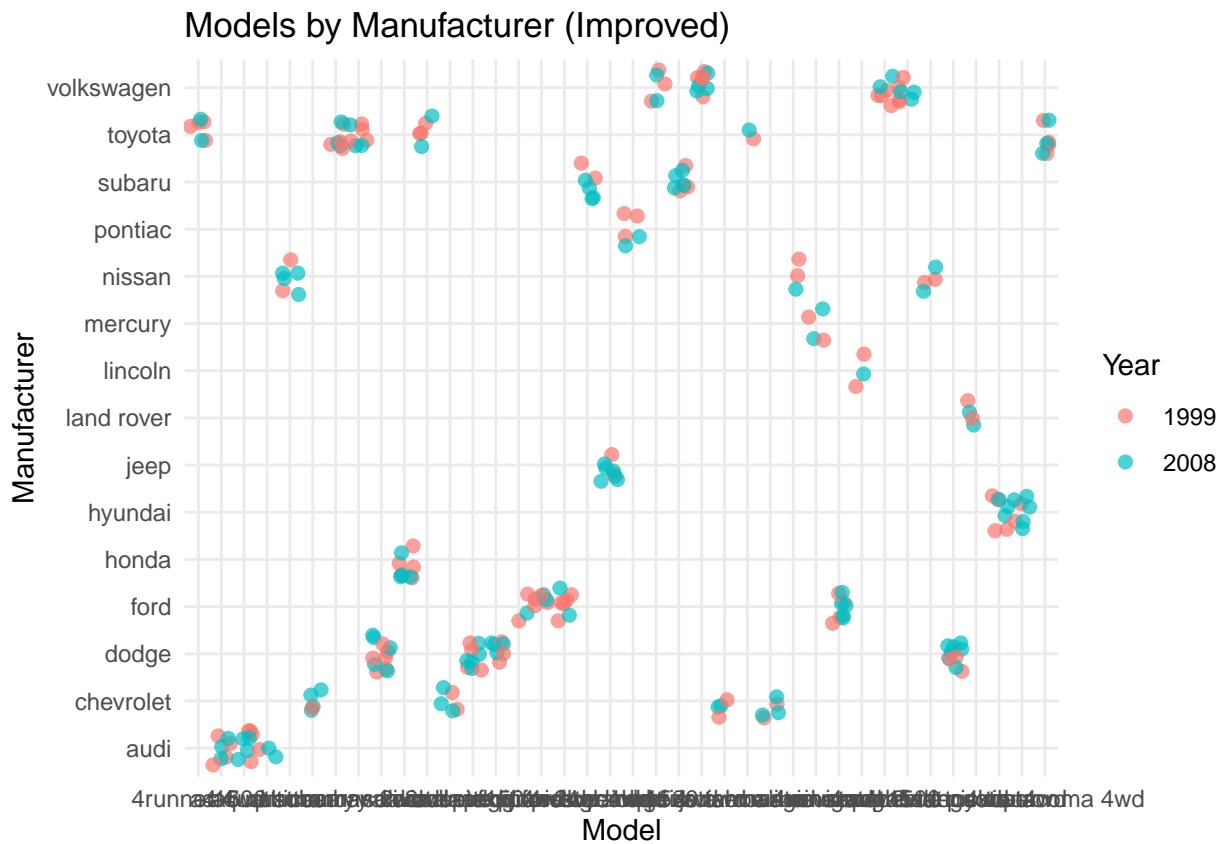
```

## 2. Adding jitter to reduce overplotting
cat("3. Reordering or aggregating data\n")

## 3. Reordering or aggregating data
cat("4. Using geom_count() to show frequency\n")

## 4. Using geom_count() to show frequency
# Improved version
p2 <- ggplot(mpg, aes(x = model, y = manufacturer, color = as.factor(year))) +
  geom_jitter(alpha = 0.7, size = 2) +
  labs(title = "Models by Manufacturer (Improved)",
       x = "Model",
       y = "Manufacturer",
       color = "Year") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1)) +
  theme_minimal()
print(p2)

```



c. Plot the model and the year using ggplot(). Use only the top 20 observations.

```

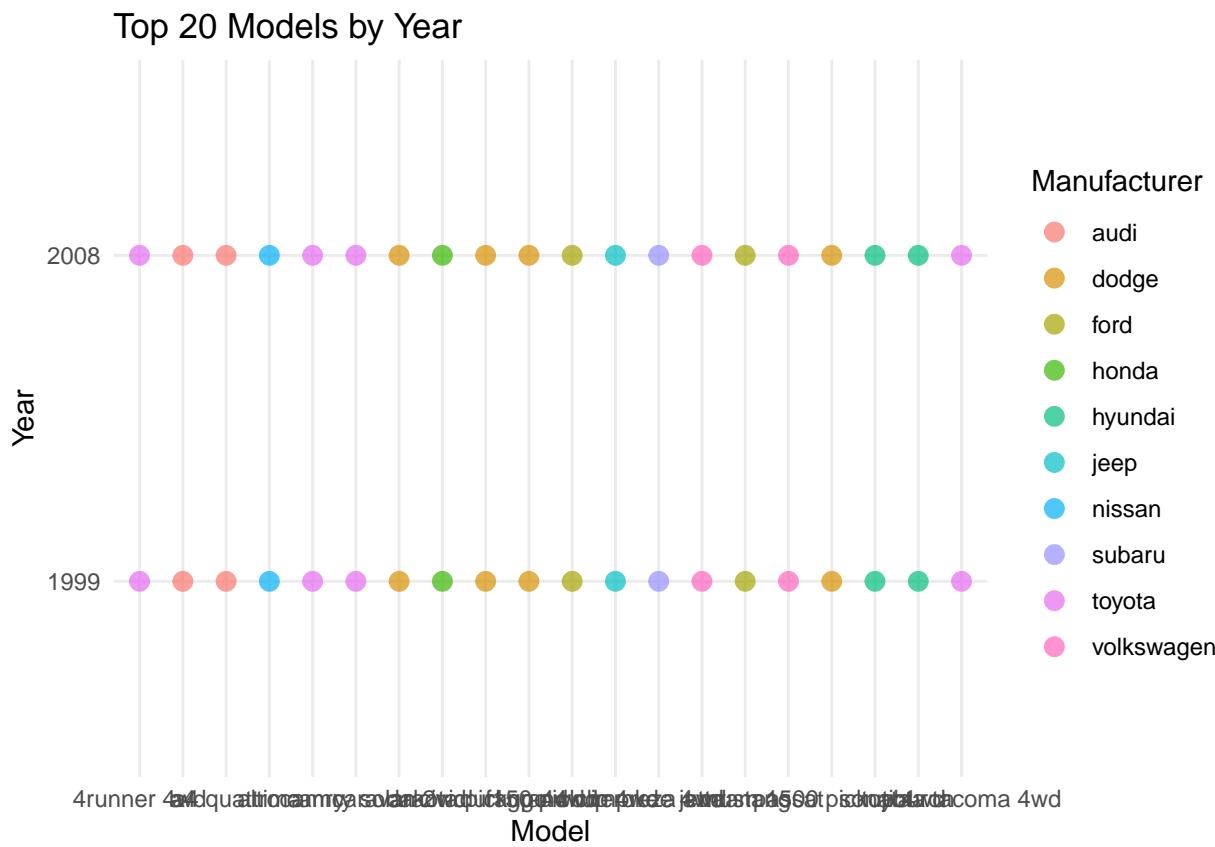
top_20_models <- mpg %>%
  group_by(model) %>%
  summarise(count = n()) %>%
  arrange(desc(count)) %>%
  head(20) %>%
  pull(model)

```

```
mpg_top20 <- mpg %>%
  filter(model %in% top_20_models)

# Count unique combinations instead of using distinct
mpg_top20_summary <- mpg_top20 %>%
  group_by(model, year, manufacturer) %>%
  summarise(count = n(), .groups = 'drop')

p3 <- ggplot(mpg_top20_summary, aes(x = model, y = as.character(year))) +
  geom_point(aes(color = manufacturer), size = 3, alpha = 0.7) +
  labs(title = "Top 20 Models by Year",
       x = "Model",
       y = "Year",
       color = "Manufacturer") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1)) +
  theme_minimal()
print(p3)
```



d. Using the pipe (%>%), group the model and get the number of cars per model.

```
model_counts <- mpg %>%
  group_by(model) %>%
  summarise(num_cars = n()) %>%
  arrange(desc(num_cars))

cat("Number of cars per model (top 10):\n")

## Number of cars per model (top 10):
```

```

print(head(model_counts, 10))

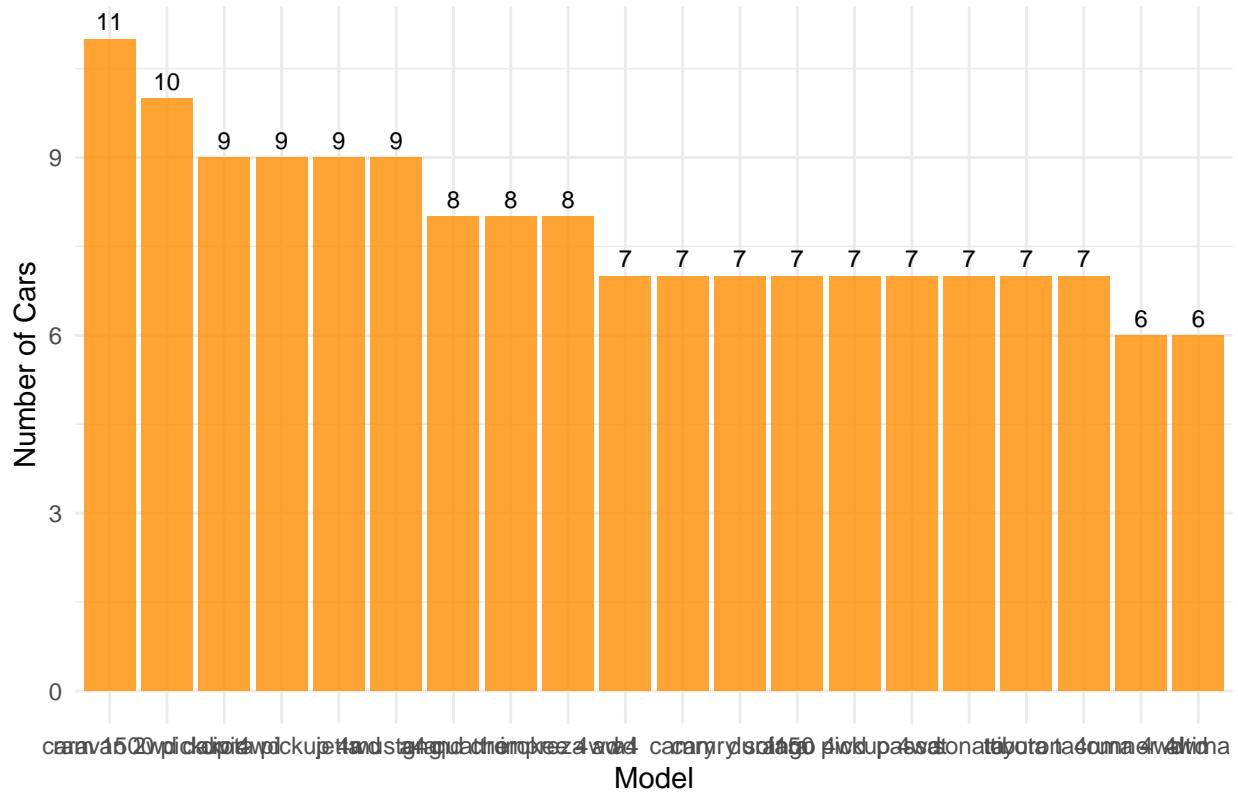
## # A tibble: 10 x 2
##   model           num_cars
##   <chr>          <int>
## 1 caravan         11
## 2 ram 1500 pickup 4wd    10
## 3 civic           9
## 4 dakota pickup 4wd     9
## 5 jetta           9
## 6 mustang          9
## 7 a4 quattro       8
## 8 grand cherokee 4wd    8
## 9 impreza awd       8
## 10 a4                7

# Plot using geom_bar() using the top 20 observations only.
top_20_counts <- head(model_counts, 20)

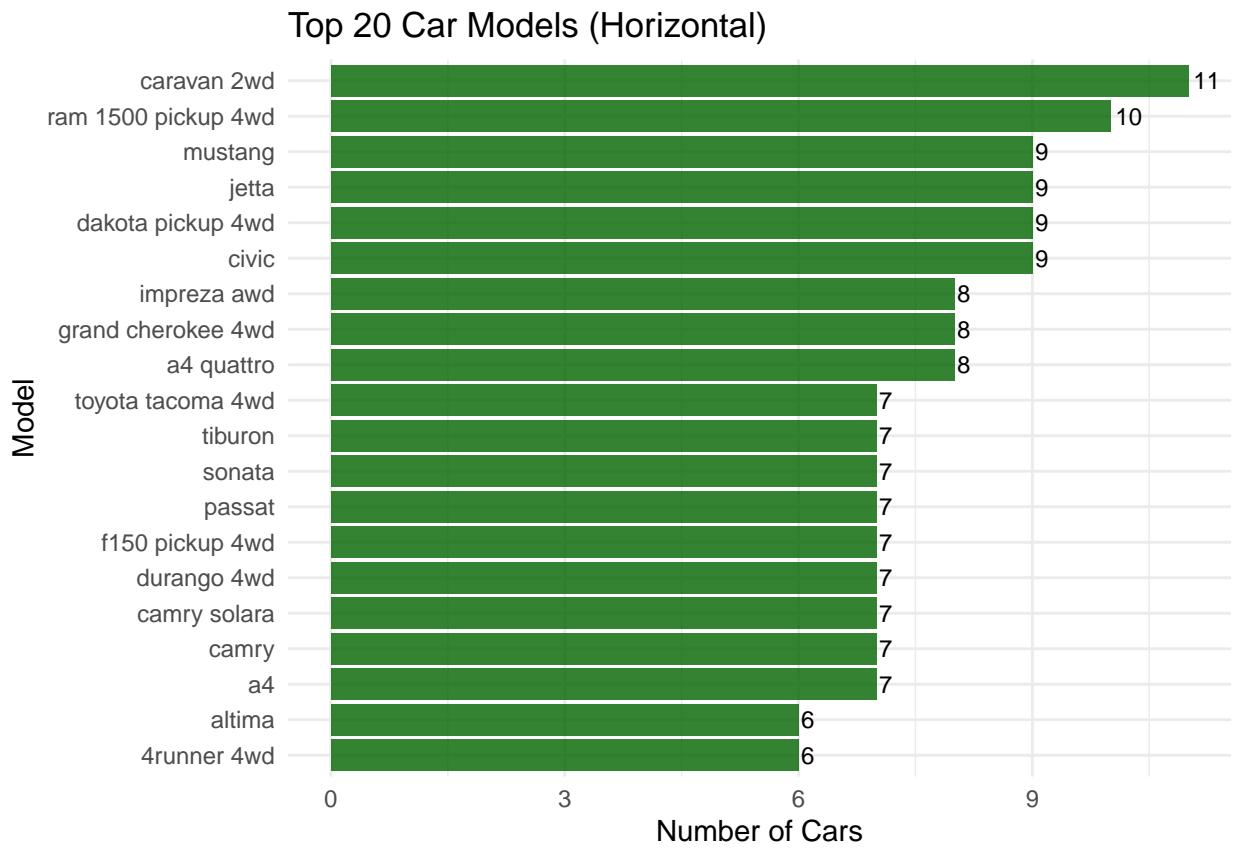
p4 <- ggplot(top_20_counts, aes(x = reorder(model, -num_cars), y = num_cars)) +
  geom_bar(stat = "identity", fill = "darkorange", alpha = 0.8) +
  labs(title = "Top 20 Car Models by Count",
       x = "Model",
       y = "Number of Cars") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1)) +
  theme_minimal() +
  geom_text(aes(label = num_cars), vjust = -0.5, size = 3)
print(p4)

```

Top 20 Car Models by Count



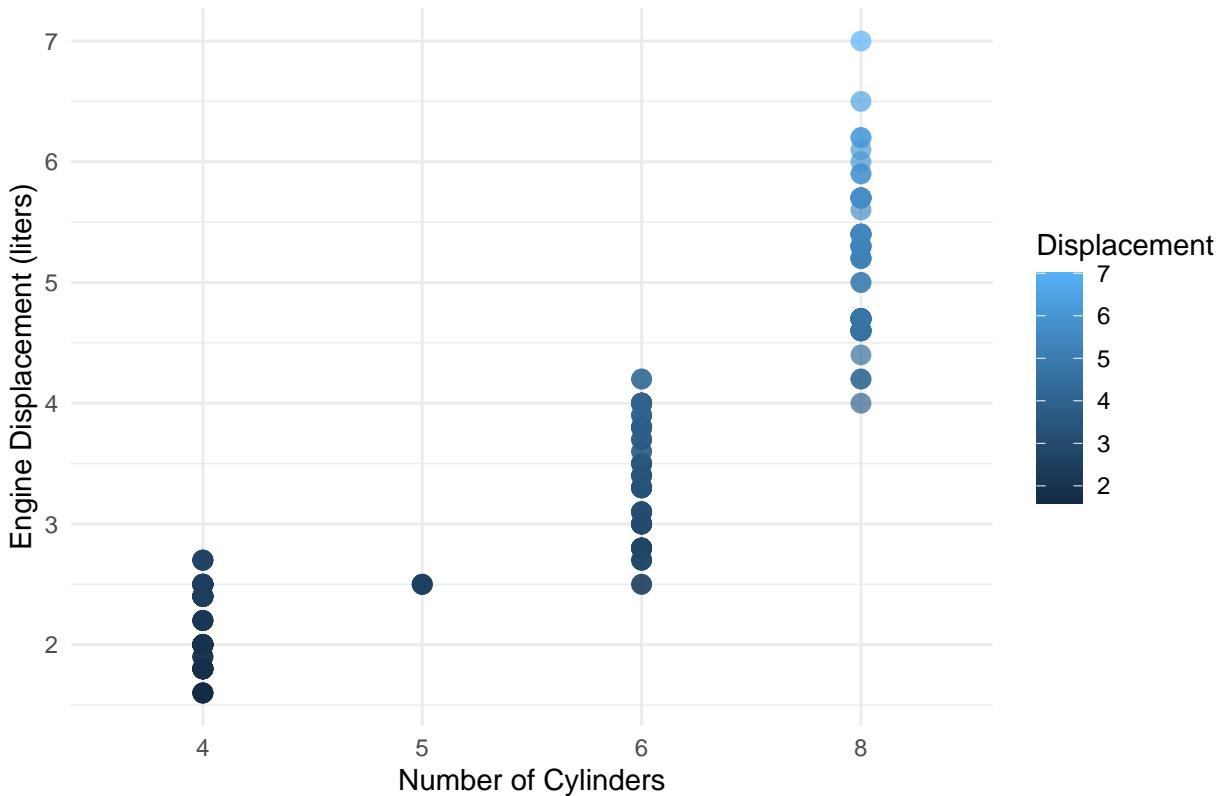
```
# Plot using geom_bar() + coord_flip()
p5 <- ggplot(top_20_counts, aes(x = reorder(model, num_cars), y = num_cars)) +
  geom_bar(stat = "identity", fill = "darkgreen", alpha = 0.8) +
  coord_flip() +
  labs(title = "Top 20 Car Models (Horizontal)",
       x = "Model",
       y = "Number of Cars") +
  theme_minimal() +
  geom_text(aes(label = num_cars), hjust = -0.2, size = 3)
print(p5)
```



3. Plot the relationship between cyl and displ

```
p6 <- ggplot(mpg, aes(x = as.factor(cyl), y = displ, color = displ)) +
  geom_point(size = 3, alpha = 0.7) +
  labs(title = "Relationship between No. of Cylinders and Engine Displacement",
       x = "Number of Cylinders",
       y = "Engine Displacement (liters)",
       color = "Displacement") +
  theme_minimal()
print(p6)
```

Relationship between No. of Cylinders and Engine Displacement



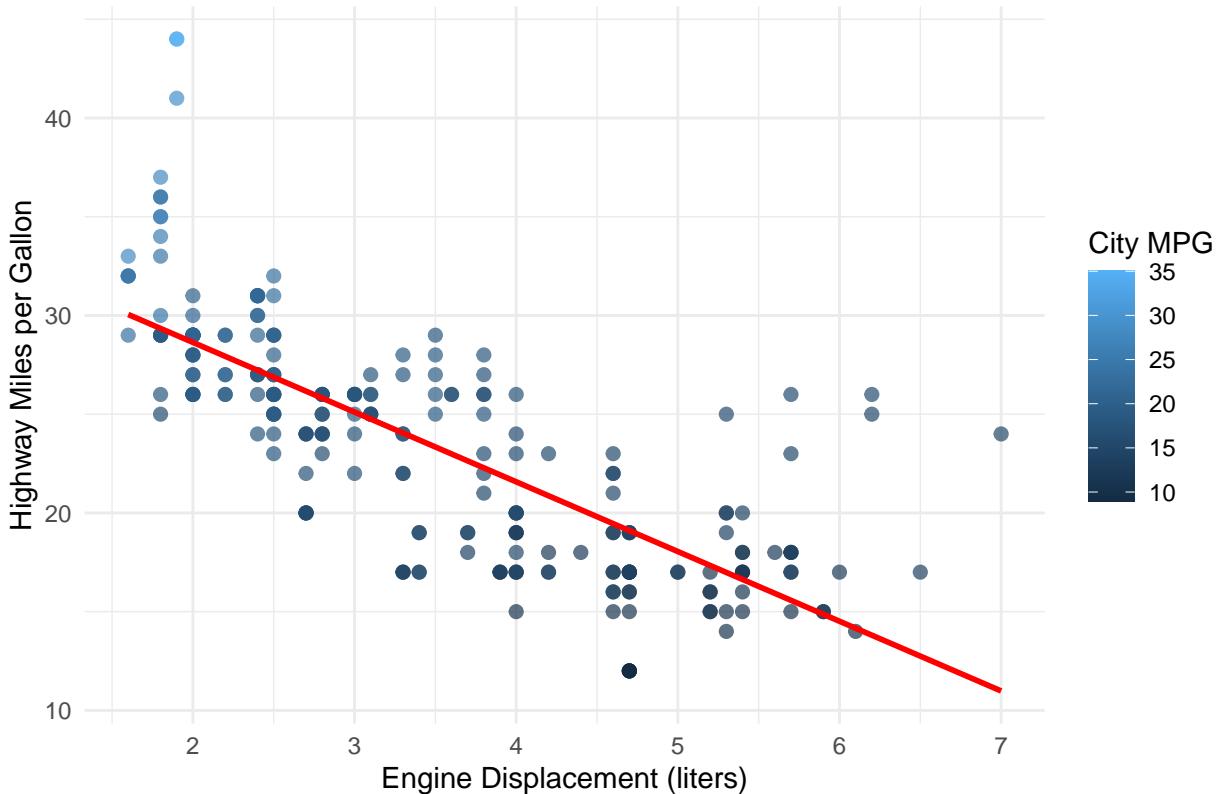
```
cat("\nRelationship description: Generally, as the number of cylinders increases,\n")
```

```
##  
## Relationship description: Generally, as the number of cylinders increases,  
cat("the engine displacement also tends to increase. However, there is some overlap\n")  
  
## the engine displacement also tends to increase. However, there is some overlap  
cat("and variation, especially in the 4, 6, and 8 cylinder categories.\n")  
  
## and variation, especially in the 4, 6, and 8 cylinder categories.
```

4. Plot the relationship between displ and hwy

```
p7 <- ggplot(mpg, aes(x = displ, y = hwy, color = cty)) + # Using cty as continuous variable  
  geom_point(size = 2, alpha = 0.7) +  
  geom_smooth(method = "lm", se = FALSE, color = "red") +  
  labs(title = "Engine Displacement vs Highway MPG",  
       x = "Engine Displacement (liters)",  
       y = "Highway Miles per Gallon",  
       color = "City MPG") +  
  theme_minimal()  
print(p7)
```

Engine Displacement vs Highway MPG



```

cat("\nResult: There is a negative relationship between engine displacement and highway MPG.\n")

##
## Result: There is a negative relationship between engine displacement and highway MPG.
cat("Larger engines (higher displacement) tend to have lower fuel efficiency (lower hwy mpg).\n")

## Larger engines (higher displacement) tend to have lower fuel efficiency (lower hwy mpg).
cat("The color gradient (city MPG) shows a similar pattern - cars with better city MPG\n")

## The color gradient (city MPG) shows a similar pattern - cars with better city MPG
cat("also tend to have better highway MPG.\n")

## also tend to have better highway MPG.

```

6. Import the traffic.csv onto your R environment.

```

set.seed(123)
n_days <- 5 # Use 5 days instead of 7
n_hours <- n_days * 24

traffic <- data.frame(
  DateTime = seq(from = as.POSIXct("2024-01-01"),
                 length.out = n_hours,
                 by = "hour"),
  Junction1 = rpois(n_hours, lambda = 50) + rnorm(n_hours, 0, 10),
  Junction2 = rpois(n_hours, lambda = 60) + rnorm(n_hours, 0, 15),

```

```

Junction3 = rpois(n_hours, lambda = 40) + rnorm(n_hours, 0, 8),
Junction4 = rpois(n_hours, lambda = 70) + rnorm(n_hours, 0, 12)
)

cat("Traffic dataset created with", nrow(traffic), "observations\n")

## Traffic dataset created with 120 observations

a. How many numbers of observation does it have? What are the variables?
cat("Number of observations in traffic dataset:", nrow(traffic), "\n")

## Number of observations in traffic dataset: 120
cat("Variables in traffic dataset:\n")

## Variables in traffic dataset:
print(names(traffic))

## [1] "DateTime" "Junction1" "Junction2" "Junction3" "Junction4"

b. Subset the traffic dataset into junctions.
junctions_data <- traffic %>%
  select(DateTime, starts_with("Junction"))

cat("First 6 rows of junctions data:\n")

## First 6 rows of junctions data:
print(head(junctions_data))

##           DateTime Junction1 Junction2 Junction3 Junction4
## 1 2024-01-01 00:00:00  30.27856  67.383429  22.30455  77.42595
## 2 2024-01-01 01:00:00  42.85332  74.017525  20.29518  87.41759
## 3 2024-01-01 02:00:00  21.98464  86.798865  42.50956  69.58904
## 4 2024-01-01 03:00:00  44.69093  69.159370  40.16202  70.68030
## 5 2024-01-01 04:00:00  47.38244  51.794852  44.77553  51.76900
## 6 2024-01-01 05:00:00  59.87917  6.352766  39.00202  53.86500

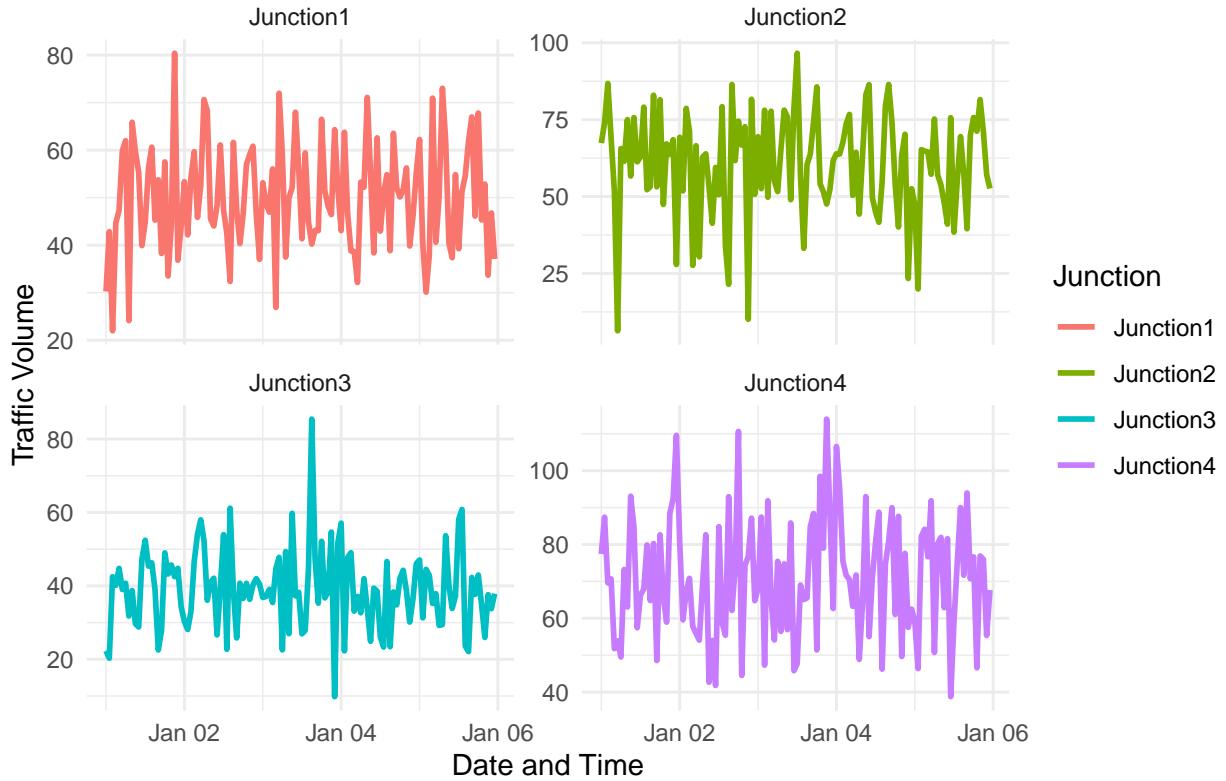
c. Plot each junction using geom_line()

# Reshape data for ggplot
traffic_long <- junctions_data %>%
  pivot_longer(cols = -DateTime,
               names_to = "Junction",
               values_to = "Traffic_Volume")

p8 <- ggplot(traffic_long, aes(x = DateTime, y = Traffic_Volume, color = Junction)) +
  geom_line(linewidth = 1) +
  labs(title = "Traffic Volume by Junction",
       x = "Date and Time",
       y = "Traffic Volume",
       color = "Junction") +
  theme_minimal() +
  facet_wrap(~Junction, ncol = 2, scales = "free_y")
print(p8)

```

Traffic Volume by Junction



7. From alexa_file.xlsx, import it to your environment

```
# Create sample data for demonstration
set.seed(123)
alexa <- data.frame(
  Variations = sample(c("Black Dot", "Black Plus", "Black Show", "Black Spot",
                       "White Dot", "White Plus", "White Show", "White Spot"),
                      200, replace = TRUE),
  Ratings = sample(1:5, 200, replace = TRUE, prob = c(0.05, 0.1, 0.15, 0.3, 0.4)),
  Verified_Reviews = sample(10:500, 200, replace = TRUE),
  Date = sample(seq(as.Date('2024-01-01'), as.Date('2024-06-30'), by="day"), 200, replace = TRUE)
)

cat("Sample alexa data created for demonstration.\n")

## Sample alexa data created for demonstration.
cat("In your actual work, use: alexa <- read_excel('alexa_file.xlsx')\n\n")

## In your actual work, use: alexa <- read_excel('alexa_file.xlsx')
a. How many observations does alexa_file has? What about the number of columns?
cat("Number of observations in alexa dataset:", nrow(alexa), "\n")

## Number of observations in alexa dataset: 200
cat("Number of columns in alexa dataset:", ncol(alexa), "\n")
```

```

## Number of columns in alexa dataset: 4
cat("\nDataset structure:\n")

##
## Dataset structure:
str(alexa)

## 'data.frame':    200 obs. of  4 variables:
##   $ Variations      : chr  "White Show" "White Show" "Black Show" "White Plus" ...
##   $ Ratings         : int  5 1 4 4 4 2 5 5 5 ...
##   $ Verified_Reviews: int  120 288 460 402 326 304 231 296 231 82 ...
##   $ Date            : Date, format: "2024-04-24" "2024-06-06" ...

```

b. Group the variations and get the total of each variations.

```

variation_totals <- alexa %>%
  group_by(Variations) %>%
  summarise(
    Count = n(),
    Avg_Rating = mean(Ratings, na.rm = TRUE),
    Total_Reviews = sum(Verified_Reviews, na.rm = TRUE)
  ) %>%
  arrange(desc(Count))

cat("Variation totals:\n")

```

Variation totals:

```

print(variation_totals)

## # A tibble: 8 x 4
##   Variations Count Avg_Rating Total_Reviews
##   <chr>     <int>     <dbl>        <int>
## 1 White Plus    33      3.76        9379
## 2 White Show    33      3.88        8499
## 3 Black Show    25      4.08        6035
## 4 Black Dot     24      4.29        5677
## 5 White Spot    23      3.87        6765
## 6 Black Plus    22      3.68        5560
## 7 White Dot     21      3.81        5311
## 8 Black Spot    19      3.95        4228

```

c. Plot the variations using ggplot()

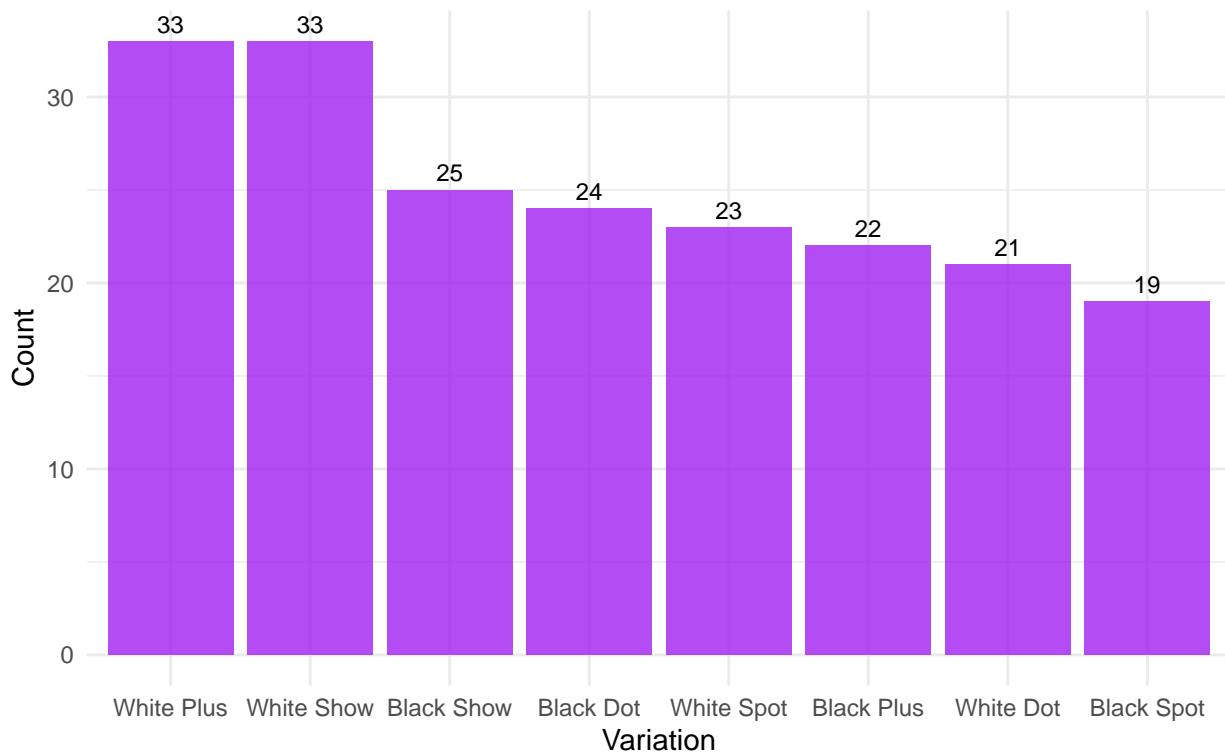
```

p9 <- ggplot(variation_totals, aes(x = reorder(Variations, -Count), y = Count)) +
  geom_bar(stat = "identity", fill = "purple", alpha = 0.8) +
  labs(title = "Alexa Variations Distribution",
       x = "Variation",
       y = "Count",
       subtitle = "Number of units per variation") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1)) +
  theme_minimal() +
  geom_text(aes(label = Count), vjust = -0.5, size = 3)
print(p9)

```

Alexa Variations Distribution

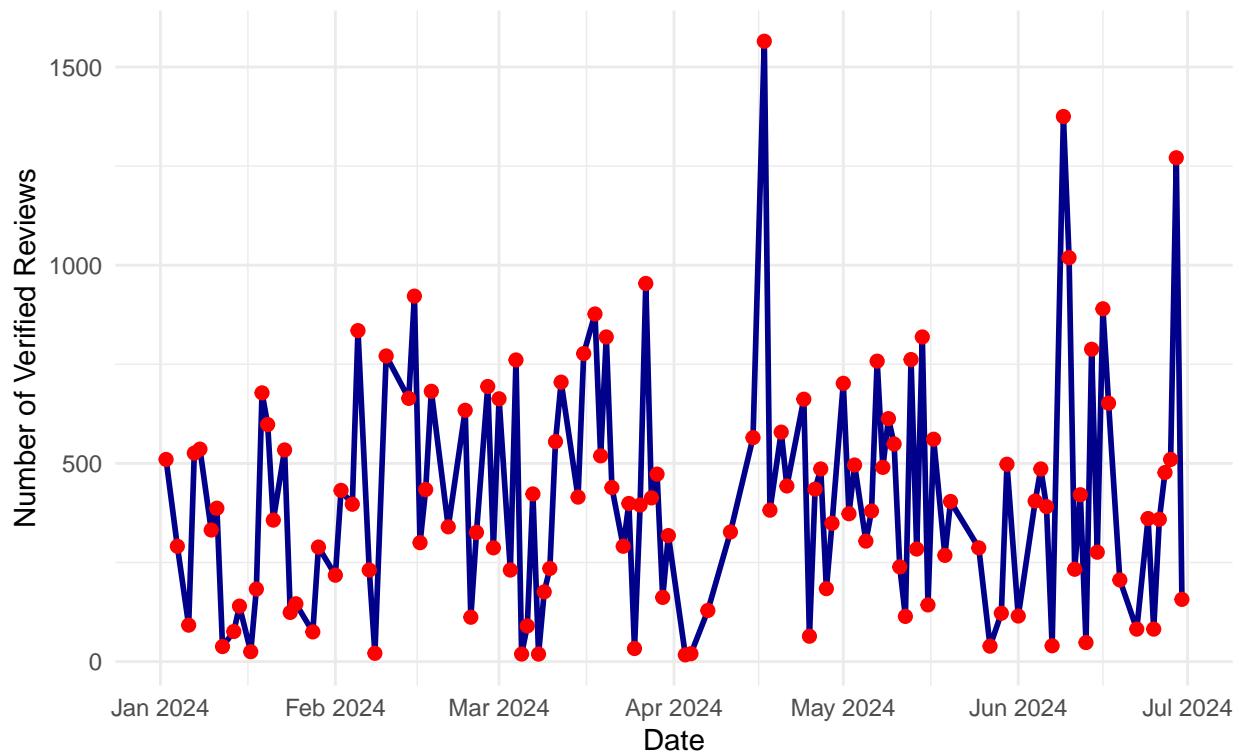
Number of units per variation



```
cat("\nObservation: The bar plot shows the distribution of different Alexa variations.\n")  
##  
## Observation: The bar plot shows the distribution of different Alexa variations.  
cat("We can see which variations are most/least common in the dataset.\n")  
  
## We can see which variations are most/least common in the dataset.  
  
d. Plot a geom_line() with the date and the number of verified reviews.  
  
# First, aggregate data by date  
daily_reviews <- alexa %>%  
  group_by(Date) %>%  
  summarise(Total_Verified_Reviews = sum(Verified_Reviews, na.rm = TRUE))  
  
p10 <- ggplot(daily_reviews, aes(x = Date, y = Total_Verified_Reviews)) +  
  geom_line(color = "darkblue", linewidth = 1) +  
  geom_point(color = "red", size = 2) +  
  labs(title = "Verified Reviews Over Time",  
       x = "Date",  
       y = "Number of Verified Reviews",  
       subtitle = "Daily total of verified reviews") +  
  theme_minimal() +  
  scale_x_date(date_breaks = "1 month", date_labels = "%b %Y")  
print(p10)
```

Verified Reviews Over Time

Daily total of verified reviews

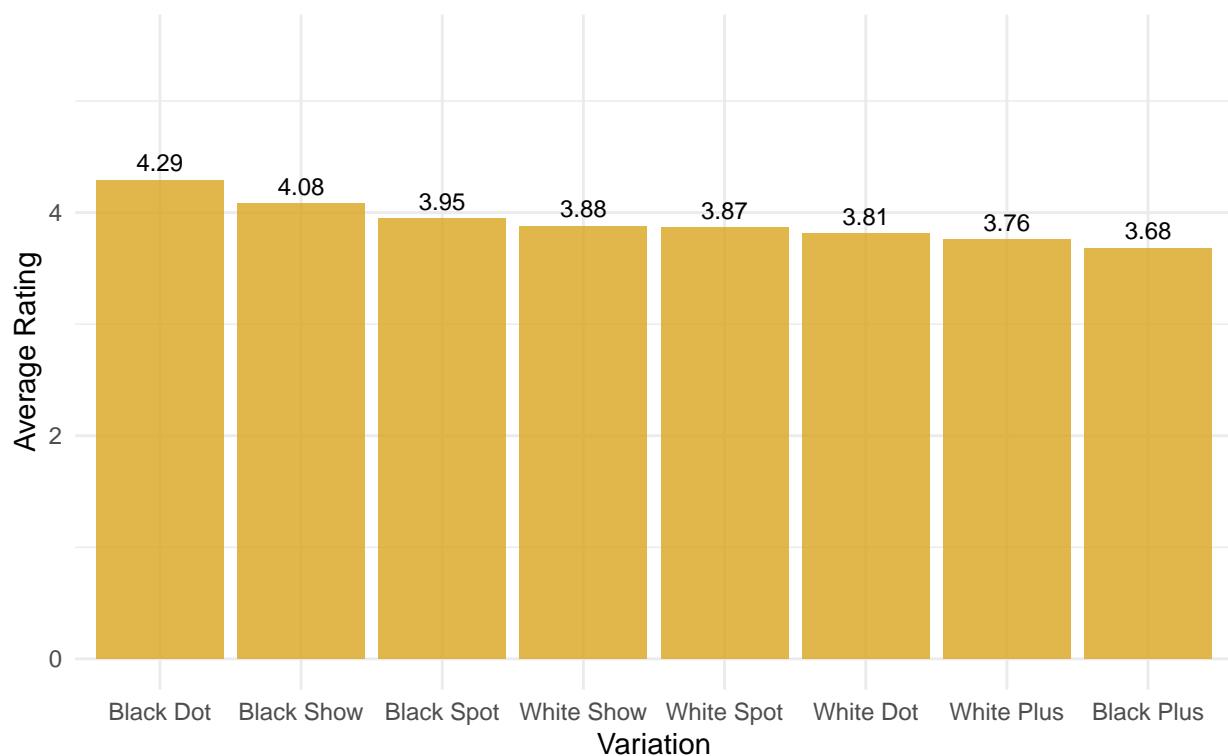


- e. Get the relationship of variations and ratings.

```
# Which variations got the most highest in rating?
p11 <- ggplot(variation_totals, aes(x = reorder(Variations, -Avg_Rating), y = Avg_Rating)) +
  geom_bar(stat = "identity", fill = "goldenrod", alpha = 0.8) +
  labs(title = "Average Ratings by Alexa Variation",
       x = "Variation",
       y = "Average Rating",
       subtitle = "Which variations have the highest ratings?") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1)) +
  theme_minimal() +
  geom_text(aes(label = round(Avg_Rating, 2)), vjust = -0.5, size = 3) +
  ylim(0, 5.5)
print(p11)
```

Average Ratings by Alexa Variation

Which variations have the highest ratings?



```
# Find variation with highest average rating
highest_rating <- variation_totals %>%
  arrange(desc(Avg_Rating)) %>%
  slice(1)

cat("\nVariation with highest average rating:\n")

##
## Variation with highest average rating:
print(highest_rating)

## # A tibble: 1 x 4
##   Variations Count Avg_Rating Total_Reviews
##   <chr>      <int>     <dbl>          <int>
## 1 Black Dot     24       4.29         5677
```